Encoder fine-tuning variant (explored)

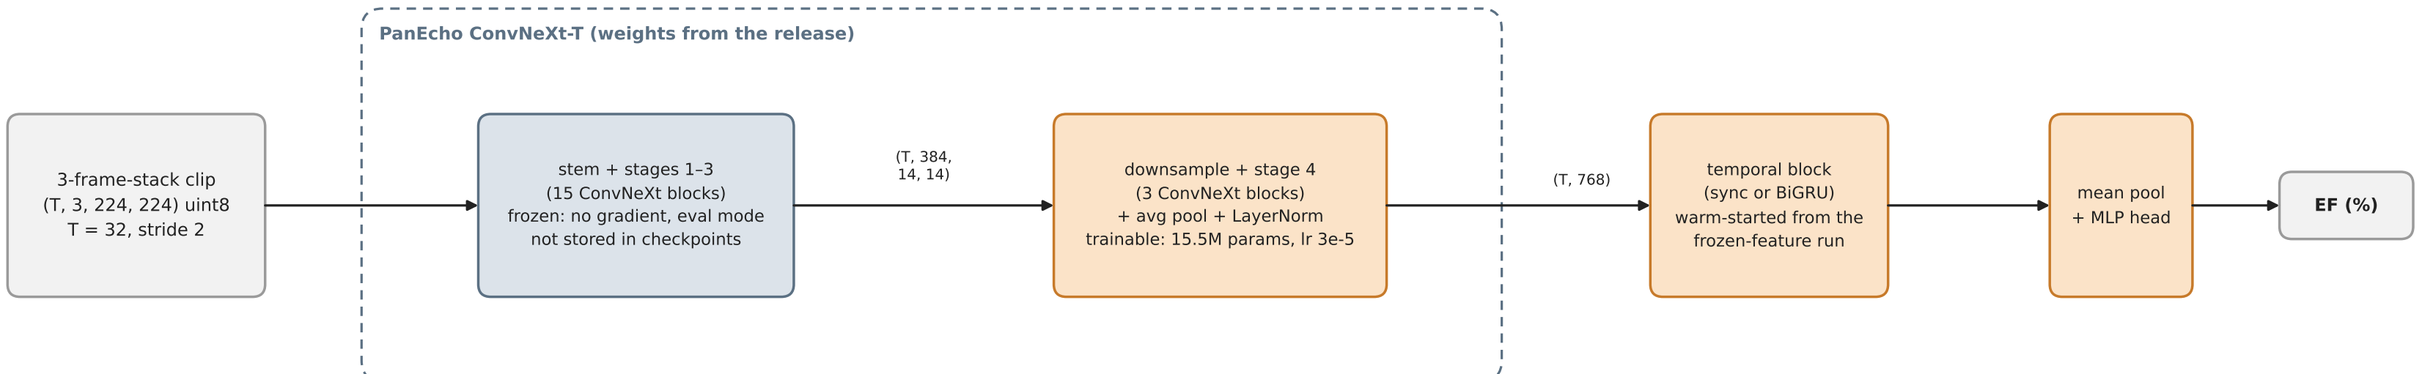

Linear probe, then fine-tune (LP-FT): head lr 1e-4, encoder lr 3e-5, 5 epochs, weight EMA 0.999. Explored and not adopted: validation error rose every epoch (the encoder overfits 7,465 training videos).
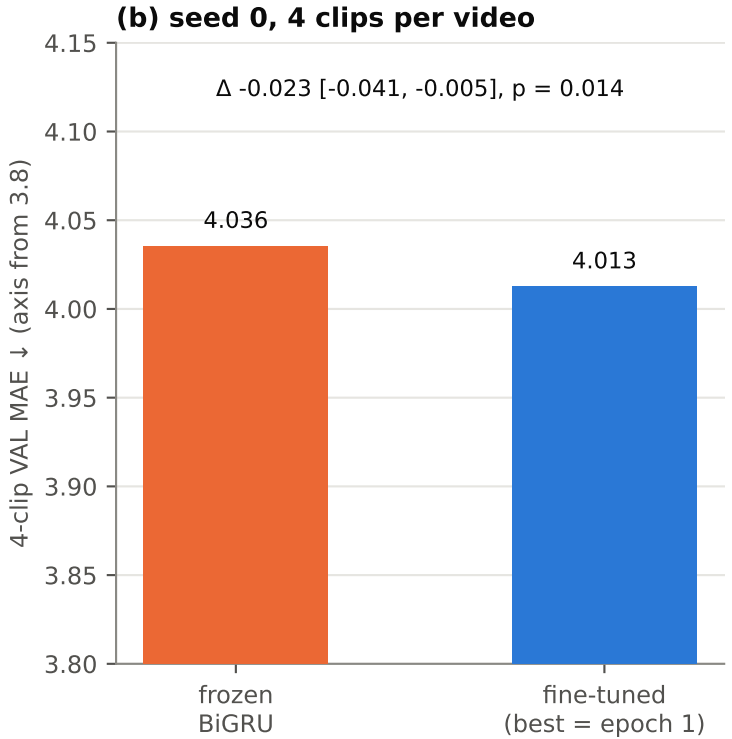
Encoder fine-tuning vs frozen BiGRU (validation set): the encoder overfits after the first epoch